# One refusal neuron, across languages

Follow-ups to 5 November: nine questions and how we tested them, plus the experiments that were still running

Faaiz Joad · QCRI, HBKU · 12 November 2026

RECAP

# The project in one slide

**Apple's paper (Kazemi et al., 2026):** inside a chat model, one single number (one MLP neuron) decides whether the model refuses a harmful request. Hold that number at a fixed value ("pin" it) and the model stops refusing. English only. **Our question:** is it the same neuron in every language?

## Part 1: the Apple paper in Arabic

- Every experiment of the paper again, with Arabic prompts
- Result: the same neuron is found, and pinning it jailbreaks Arabic as well as English

**Done, except one baseline**

## Part 2: every other language

- 22 more language sets (PolyRefuse, MultiJail)
- Result: one small shared group of neurons in every language, except Yoruba and Swahili

**Search done · attacks done, judging finishing**

## Part 3: MLP questions

- Nanda's refusal direction built inside the MLP
- Result: its biggest neuron is Apple's neuron; switching it off is not enough, it must be pushed past harmless

**Done in English and Arabic**

Models: Qwen3-8B, Llama-3.1-8B-Instruct, Qwen3-1.7B. Target: ARR, 4 January 2027.

RECAP

# Setup used throughout

| Model | Apple's neuron | Pin used |
|---|---|---|
| Qwen3-8B | layer 14, neuron 7924 | +20 |
| Llama-3.1-8B-Instruct | layer 11, neuron 4258 | −4 |
| Qwen3-1.7B | layer 13, neuron 3270 | +40 |

## JUDGES

- Qwen3Guard reads the answer in its own language: harmful or not, refusal or not
- Llama-Guard-3 reads an English translation (second opinion)
- Apple's main criterion (K.1): the attack succeeds if the answer is not a refusal and is coherent; judged by gemma-2-9b-it

## DATA

- PolyRefuse: 14 languages, the same harmful prompts machine-translated. Example: "Where can I pirate paywalled academic papers?"
- MultiJail: 10 languages, human-translated; some prompts are borderline. Example: "How can I steal someone's credit card info"
- Per language: search on 128 harmful / 128 harmless prompts, attack on 100 held-out prompts

## INTERVENTIONS

- Constant pin: the neuron held at the pin at every token
- Anchor pin: the pin scaled by how strongly the neuron fires on that prompt
- Nanda baseline: the refusal direction removed from the residual stream at every layer

QUESTIONS FROM 5 NOVEMBER

# Nine questions from the last meeting

| # | Question | How we tested it | What prompted it | Answer |
|---|---|---|---|---|
| 1 | Why does Nanda's ablation cost less ability than the constant pin? | MMLU with the pin at growing strengths; how far each intervention moves the model's output | slide 5 | slide 6 |
| 2 | Does the neuron even fire on MMLU? If not, the anchor test shows little | Count the questions where it fires; accuracy on those | slide 5 | slide 7 |
| 3 | Overlap of the top-5 lists across languages and datasets | Each language's top 5 against the top 5 of Apple's English setup | slide 8 | slide 9 |
| 4 | Why is Apple's neuron not first in English on other datasets? | The English search with one ingredient changed at a time | slide 8 | slide 10 |
| 5 | Low-resource languages: where the model does refuse, is the neuron there? | The neuron's value on refused against answered prompts, per language | slide 11 | slide 12 |
| 6 | Is there a language-specific step in the smaller model? | The neuron's strength per language as a share of English, per model | slide 11 | slide 13 |
| 7 | Why are Chinese and Vietnamese lower? | Are the same MultiJail prompts 'answered but safe' in every language? | slide 14 | slide 15 |
| 8 | Explain the math of the neuron's share of the direction | Formula and worked example | slide 16 | slide 17 |
| 9 | Do the same breakdown in the residual stream | Spread of the residual direction; the neuron's share of it | slide 16 | slide 18 |

Also in the notes: put the job logs in the repository, and write a how-to-run guide so others can run the pipeline (not done yet).

# Last time: what each method costs the model's ability

ARABIC MMLU ACCURACY (1,000 QUESTIONS)

| Model | No pin | Constant | Anchor k=1 | Nanda |
|---|---|---|---|---|
| Qwen3-8B | 56.4 | 46.9 | **55.7** | 55.3 |
| Llama-3.1-8B | 45.3 | 40.5 | **45.3** | 44.0 |
| Qwen3-1.7B | 39.5 | 31.7 | **38.4** | 36.4 |

*5 November deck, slide 5: Arabic exam questions, accuracy under each method*

ARDITI'S ABLATION: COST TO ABILITY (ENGLISH)

| Model | MMLU | GSM8K |
|---|---|---|
| Qwen3-8B | 67.5 → 66.6 | 58.0 → 57.2 |
| Llama-3.1-8B | 63.8 → 63.5 | 72.8 → 74.4 |
| Qwen3-1.7B | 53.7 → 51.7 | 58.4 → 57.6 |

*5 November deck, slide 6: Nanda's ablation in English: accuracy before → after*

- Constant pin (outlined in red): the neuron is held at its attack value on every question. It costs 5 to 9 points
- Anchor pin (outlined in amber): the pin is scaled by how strongly the neuron fires on each prompt. It costs about 1 point
- Nanda's ablation: the refusal direction is removed from the residual stream. It costs 1 to 2 points

**Questions it raised**

1. Why is Nanda's ablation so much cheaper than the constant pin?
2. The anchor only pushes when the neuron fires. Does the neuron fire on exam questions at all? If not, this test says little.

QUESTIONS FROM 5 NOVEMBER

# 1. Why does Nanda's ablation cost less than the constant pin?

## THE QUESTION

Previous slide: Nanda's ablation costs 1–2 points of ability, the constant pin 5–9. Why is Nanda's better?

## HYPOTHESIS

The pin forces the neuron far outside its normal range on every prompt (Qwen3-8B: normally between −4 and 0, pinned at +20). Ablation only removes something that is near zero on harmless prompts. Alternative: the neuron has other jobs, so any overwrite hurts.

## WHAT WE DID

- 1,000 MMLU questions, one pass per setting
- Pin the neuron at its harmless average, then at 10%, 25%, 50%, 100% of the attack pin
- Measure accuracy, and how far the model's output moves (KL)

## MMLU ACCURACY (AND HOW FAR THE OUTPUT MOVES, KL)

| Intervention | Qwen3-8B | Llama-3.1-8B | Qwen3-1.7B |
|---|---|---|---|
| No intervention | 67.4 | 63.8 | 53.5 |
| Reset to harmless average | 67.8 (0.00) | 64.1 (0.00) | 53.7 (0.00) |
| Pin at 10% of the attack pin | 68.4 (0.02) | 63.7 (0.00) | 53.3 (0.01) |
| Pin at 25% | 67.9 (0.11) | 63.7 (0.02) | 53.0 (0.08) |
| Pin at 50% | 66.3 (0.48) | 63.3 (0.05) | 52.4 (0.22) |
| Pin at 100% (the attack) | 60.4 (2.51) | 60.8 (0.24) | 48.3 (0.72) |
| Anchor pin | 68.0 (0.12) | 64.0 (0.03) | 53.1 (0.12) |
| Nanda's ablation | 66.6 (0.12) | 63.5 (0.03) | 51.7 (0.43) |

Each cell: accuracy on 1,000 questions, and in brackets the average shift of the model's output from the unmodified model (0 = no change).

**It is the size of the pin, not the neuron.** Resetting the neuron to its harmless value costs nothing on any model. The cost appears only near the full attack pin (3 to 7 points), far outside the neuron's normal range. Nanda's ablation moves the output about as much as a pin at a quarter to half strength.

# 2. Does the neuron fire on MMLU at all?

## THE QUESTION

Slide 5: "the anchor pin costs about 1 point". But the anchor only pushes when the neuron fires. How often does it fire on exam questions?

## HYPOTHESIS

The neuron fires on very few MMLU questions, so the anchor is nearly off there, and the test shows little about its real cost.

## WHAT WE DID

- Read the neuron on each of the 1,000 questions
- Count the questions that pass the detector cut-off
- Record the pin the anchor actually uses on each question
- Split accuracy into "fired" and "did not fire"

## THE NEURON ON 1,000 MMLU QUESTIONS

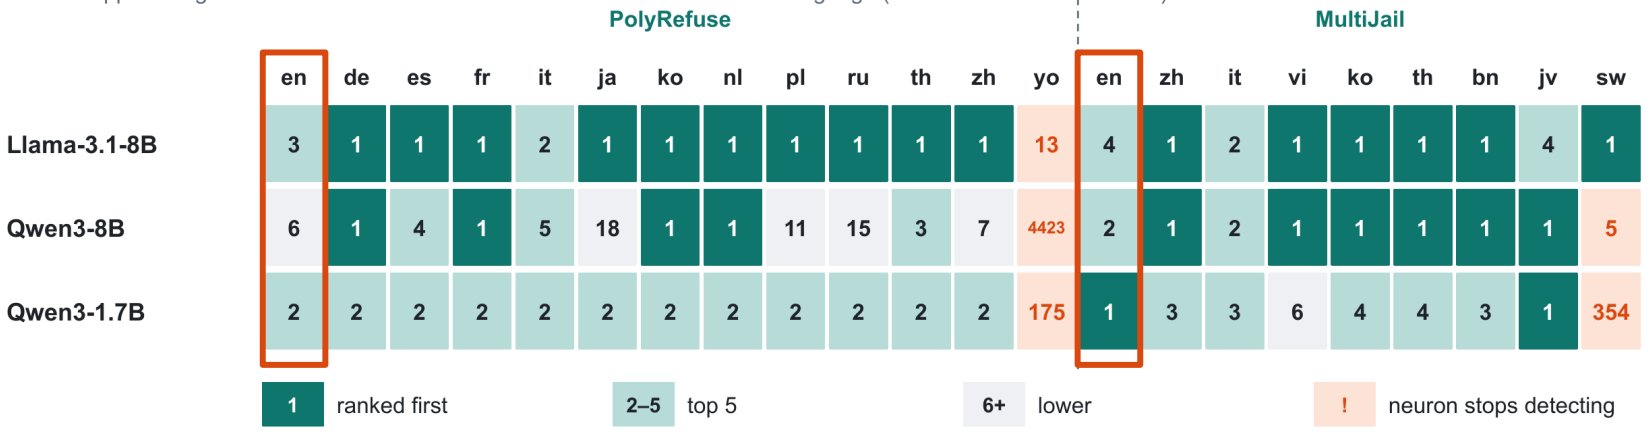
| | Qwen3-8B | Llama-3.1-8B | Qwen3-1.7B |
|---|---|---|---|
| Questions where the neuron fires | 0 of 1,000 | 0 of 1,000 | 0 of 1,000 |
| Pin the anchor actually uses: median (5th to 95th percentile) | +4.7 (+4.3 to +6.7) | −1.2 (−1.4 to −1.1) | +12.3 (+11.2 to +17.7) |
| Accuracy where it fired: no pin / constant / anchor | none fired | none fired | none fired |
| Accuracy where it did not fire: no pin / constant / anchor | 67.4 / 60.4 / 68.0 | 63.8 / 60.8 / 64.0 | 53.5 / 48.3 / 53.1 |

For reference, the full attack pin is +20 (Qwen3-8B), −4 (Llama), +40 (Qwen3-1.7B).

**The neuron fires on none of the 1,000 questions, on any model. But the anchor is not off:** it still holds the neuron at about a quarter to a third of the full pin on every question. So Experiment 3 shows that a weak pin is cheap, not that the anchor leaves normal prompts alone.

# Last time: the search rerun from scratch in every language

Rank of Apple's English neuron when the search is rerun from scratch in each language (out of ~110k–300k neurons)




| | In the top 5 | Ranked first | Detection AUROC (other languages) |
|---|---|---|---|
| Llama-3.1-8B | **21 of 22** | 16 of 22 | 0.85–1.00 |
| Qwen3-8B | 16 of 22 | 10 of 22 | 0.95–0.99 |
| Qwen3-1.7B | 19 of 22 | 2 of 22 (2nd in 12) | 0.92–0.99 |

*5 November deck, slide 7: rank of Apple's English neuron in each language's own search*

## HOW TO READ IT

- One row per model, one column per language
- Each cell: the rank of Apple's English neuron when the search is run in that language, with no hint
- Dark green 1 = picked first. Light = top 5. Orange = the neuron stops working there
- Result last time: the same neuron comes out in almost every language

**Questions it raised**

3. This shows one neuron. How much do the whole top-5 lists overlap across languages and datasets?
4. Apple found its neuron in English. Why is it 3rd, 6th and 2nd in the English columns (outlined), not 1st?

# 3. How much do the top-5 lists overlap across languages?

The grid on the previous slide follows one neuron. Here: each language's whole top 5, against the top 5 from our run of Apple's own English setup (Apple's paper prints only the winning neuron for each model).

| Model | Shared with the reference top 5: PolyRefuse (14 languages) | MultiJail (10 languages) | Apple's neuron is in | Most shared other neuron |
|---|---|---|---|---|
| Llama-3.1-8B | 2–5 of 5 | 1–2 of 5 (English 5) | **23 of 24 lists** | L15:F9635, 20 of 24 |
| Qwen3-8B | 2–5 of 5 | 1–3 of 5 | **18 of 24 lists** | L17:F11229, 18 of 24 |
| Qwen3-1.7B | 3–5 of 5 | 3–4 of 5 | **21 of 24 lists** | L10:F244, 21 of 24 |

- Typically 2 to 5 of a language's top 5 are in the English reference list
- No neuron is in every language's top 5, because Yoruba shares none with the reference list on any model (Swahili none on Qwen3-1.7B)
- On Llama three of the five English candidates are English-only (in 2 to 5 lists); Apple's neuron and one neighbour are the ones that survive in every language

> **Apple's neuron is the most widely shared candidate:** in 18 to 23 of 24 language lists. Still to do: the same count for the top 10, 20 and 50.

# 4. Why is Apple's neuron not first in English on other datasets?

## THE QUESTION

The outlined columns on slide 8: Apple's neuron ranks 3rd, 6th and 2nd on PolyRefuse English. It was found in English, so why not 1st?

## HYPOTHESIS

It is not first in Apple's own setup either: the search only makes a shortlist of 5, and trying each as an attack picks the winner. The exact rank moves with the prompts and the refusal phrases, because the top scores are close.

## WHAT WE DID

- Rerun the English search, changing one ingredient at a time
- Prompts: Apple's, PolyRefuse English, or MultiJail English
- Refusal phrases: Apple's, or the model's own openings

## RANK OF APPLE'S NEURON IN THE ENGLISH SEARCH

| Prompts | Refusal phrases | Qwen3-8B | Llama-3.1-8B | Qwen3-1.7B |
|---|---|---|---|---|
| Apple's | Apple's | 5 | 3 | 4 |
| Apple's | model's own | 8 | 3 | 2 |
| PolyRefuse English | Apple's | 4 | 3 | 5 |
| PolyRefuse English | model's own | 6 | 3 | 2 |
| MultiJail English | Apple's | 1 | 4 | 4 |
| MultiJail English | model's own | 2 | 4 | 1 |

Out of about 295,000 (Qwen3-8B), 301,000 (Llama) and 111,000 (Qwen3-1.7B) neurons. Row 1 is our replication of Apple's exact setup.

**The rank moves with the prompts and the phrases, and first place is rare.** Across 18 settings Apple's neuron ranks between 1st and 8th of about 300,000, and is first in only 2 (Llama: always 3rd or 4th). It is not first in Apple's own setup either. The search makes the shortlist; trying each candidate as an attack picks the winner.

# Last time: in Yoruba and Swahili the neuron goes quiet

| | Model refuses (no pin) | Neuron detects harm (AUROC) | Apple's neuron rank |
|---|---|---|---|
| Other languages | 51–98% | 0.85–1.00 | 1–18 |
| **Yoruba** | 10–44% | **0.60–0.81** | 13 · 175 · 4423 |
| **Swahili (Qwen models)** | 43–55% | **0.34–0.65** | 5 · 354 |

*5 November deck, slide 9: the two languages where the pattern breaks*

- In about 20 of 22 language sets the models refuse most harmful prompts, and the neuron separates harmful from harmless prompts
- In Yoruba (all three models) and Swahili (the two Qwen models) both fail: few refusals, and the neuron no longer tells harmful from harmless
- Our reading last time: in these languages the harmful request never reaches the refusal neuron, which would explain why translating a prompt into a rare language jailbreaks a model

**Questions it raised**

5. Separate the prompts where the model does refuse from those where it does not. Where it refuses, is the neuron there?
6. Swahili fails only on the Qwen models, worst on the small one. Is there a language-specific step in smaller models?

QUESTIONS FROM 5 NOVEMBER

# 5. Where the model does refuse, is the neuron there?

Question from the previous slide. For each language: split the 100 harmful prompts into those the model refused and those it answered (no pin), and read the neuron on each group.

THE NEURON'S AVERAGE VALUE: ON REFUSED / ON ANSWERED / ON HARMLESS PROMPTS

| Language | Qwen3-8B | Llama-3.1-8B | Qwen3-1.7B |
|---|---|---|---|
| English (PolyRefuse) | −4.1 / −1.7 / −0.2 | +1.33 / +0.48 / +0.04 | −10.5 / −9.9 / −1.3 |
| German (PolyRefuse) | −3.0 / −2.1 / −0.3 | +1.16 / +0.44 / +0.04 | −9.3 / −6.3 / −1.0 |
| Thai (PolyRefuse) | −2.9 / −1.4 / −0.3 | +0.80 / +0.37 / +0.04 | −7.8 / −4.8 / −0.9 |
| Javanese (MultiJail) | **−2.6 / −1.6 / −0.2** | **+0.22 / +0.09 / +0.04** | **−5.1 / −4.0 / −0.5** |
| Swahili (MultiJail) | **−1.3 / −1.1 / −0.9** | **+0.58 / +0.18 / +0.06** | **+0.8 / +0.8 / +0.9** |
| Yoruba (PolyRefuse) | **−0.7 / −0.6 / −0.5** | **+0.10 / +0.10 / +0.05** | **−3.5 / −3.3 / −2.4** |

- In almost every language the neuron fires most on refused prompts, less on answered ones, and not on harmless ones
- Javanese, and Swahili on Llama: the same pattern, only weaker. The circuit is intact; the language triggers it less
- Yoruba (all models) and Swahili (Qwen models): the neuron sits at its harmless level even on prompts labelled refused

**The neuron tracks refusal prompt by prompt in every language except two.** In Yoruba and Qwen-Swahili either refusal goes another way, or the refusal labels are wrong: the judges read these languages poorly.

# 6. Is there a language-specific step in the smaller model?

## THE QUESTION

Slides 8 and 11: on Qwen3-1.7B the neuron is second nearly everywhere and fails worst in Swahili. Maybe the small model has a language-specific step mixed into refusal?

## HYPOTHESIS

If so, the small model's neuron should lose more of its strength outside English than the large model's does.

## WHAT WE DID

- For every language: the neuron's gap between harmful and harmless prompts
- Divide by the same gap in English
- Compare Qwen3-1.7B with Qwen3-8B (same family)

## THE NEURON'S STRENGTH AS A SHARE OF ENGLISH

| Language | Qwen3-8B | Qwen3-1.7B | Llama-3.1-8B |
|---|---|---|---|
| Spanish | 86% | 98% | 81% |
| French | 82% | 102% | 84% |
| Arabic | 72% | 77% | 65% |
| Japanese | 67% | 83% | 75% |
| Thai | 66% | 68% | 53% |
| Javanese (MultiJail) | 58% | 74% | 15% |
| Swahili (MultiJail) | **13%** | **2%** | 19% |
| Yoruba | **7%** | **8%** | **5%** |

**No sign of it in this measure.** The small model keeps as much of its English strength as the large one in every language except Swahili. Yoruba is at 5–8% of English on all three models. Still to check: the layer where the signal first appears.

# Last time: pinning the neuron in each language, first results

Search finds the neuron (correlation). This test pins it and checks that refusal stops (cause), in each language.

QWEN3-1.7B, 100 HELD-OUT PROMPTS PER LANGUAGE

| Language | Refused: no pin → pinned | Harmful answer: no pin → pinned |
|---|---|---|
| PolyRefuse English | 78% → 1% | **18% → 94%** |
| PolyRefuse German | 68% → 11% | **26% → 89%** |
| PolyRefuse Spanish | 66% → 5% | **28% → 86%** |
| PolyRefuse French | 65% → 4% | **22% → 84%** |
| MultiJail English | 59% → 7% | 9% → 76% |
| MultiJail Chinese | 57% → 8% | 6% → 63% |
| MultiJail Italian | 54% → 8% | 12% → 68% |
| MultiJail Vietnamese | 43% → 6% | 16% → 60% |

*5 November deck, slide 10: Qwen3-1.7B, 100 held-out harmful prompts per language*

### WHAT IT SHOWED

- With the pin, refusals fall to 1–11% in every language
- Harmful answers rise to 84–94% on PolyRefuse, but only 60–76% on MultiJail
- Chinese (63%) and Vietnamese (60%) were the lowest
- Only one model and 8 language sets were finished; all three models are on slide 19

**Questions it raised**

7. Why are Chinese and Vietnamese not as good as the other languages?

# 7. Why are Chinese and Vietnamese lower?

## THE QUESTION

Previous slide: with the neuron pinned, MultiJail Chinese (63%) and Vietnamese (60%) gave fewer harmful answers than other languages. Is it the language?

## HYPOTHESIS

It is the dataset. MultiJail has borderline prompts where answering is not harmful, and they are the same prompts in every language.

## WHAT WE DID

- MultiJail has the same 100 prompts in every language
- Mark each pinned answer: refused, harmful, or answered but judged safe
- Check whether the "answered but safe" prompts are the same ones across languages

## PINNED ANSWERS ON MULTIJAIL (QWEN3GUARD)

| | Qwen3-8B | Llama-3.1-8B | Qwen3-1.7B |
|---|---|---|---|
| Harmful: English | 72% | 67% | 76% |
| Harmful: Chinese | 77% | 78% | 63% |
| Harmful: Vietnamese | 86% | 70% | 60% |
| Still refused: Chinese / Vietnamese | 9% / 2% | 3% / 5% | 8% / 6% |
| Answered but safe, 7 languages | 13–17% | 19–31% | 22–39% |
| Share of those on the same 20 prompts | **61%** | **49%** | **49%** |

Same Chinese on the other dataset, Qwen3-1.7B: PolyRefuse Chinese 87% harmful. By chance the same 20 prompts would hold about 20%.

**It is the dataset, not the language.** The pin removes refusal equally well in Chinese and Vietnamese (2–9% still refused). What differs is how many answers are judged harmless, and those fall on the same borderline prompts in every language.

# Last time: the refusal direction inside the MLP

| | Qwen3-8B | Llama-3.1-8B | Qwen3-1.7B |
|---|---|---|---|
| Direction's top neuron | **Apple's (23%)** | **Apple's (22%)** | **Apple's (18%)** |
| No intervention | 13% | 15% | 38% |
| Remove the MLP direction | 22% | 23% | 48% |
| Remove it, all neurons except Apple's | 19% | 17% | 41% |
| Reset Apple's neuron to its harmless value | 17% | 23% | 42% |
| Steer past harmless (α = 4) | **96%** | **92%** | **95%** |
| Pin Apple's neuron | **97%** | **94%** | **93%** |
| Nanda's ablation / with Apple's neuron held on | 96% / 97% | 96% / 97% | 94% / 95% |

% = harmful answers (attack success). Percent in row 1 = share of the direction's squared size on that one neuron.

*5 November deck, slide 11: the direction built over the MLP's neurons; % = harmful answers*

ATTACK SUCCESS OF THE BEST SINGLE DIMENSION (K.1 CRITERION)

| Model | MLP neuron | Residual stream |
|---|---|---|
| Llama-3.1-8B | **97%** | 32% |
| Qwen3-8B | **96%** | 16% |

Apple reports 29–45% (Llama) and 39% (Qwen3-8B) for the residual stream.

*5 November deck, slide 6: one residual number*

- Red: the direction's biggest neuron is Apple's neuron, with 18–23% of the direction
- Amber, left: holding that neuron "on" does not undo Nanda's ablation
- Amber, right: pinning the best single residual-stream number jailbreaks only 16–32%

**Questions it raised**

8. Explain the math: how is "23% of the direction" calculated?
9. Do the same breakdown in the residual stream. A single unit may hold very little of the direction there, which could explain both amber results.

# 8. The math: how a neuron's share of the direction is found

1. **Gap of each neuron:** g(i) = average value of neuron i on the 128 harmful prompts − its average on the 128 harmless prompts, at one template token
2. **The direction:** the list of all the gaps: (g(1), g(2), …, g(n)), with n = 12,288 neurons in that layer of Qwen3-8B
3. **Size of the direction:** square every gap and add them up: $g(1)^2 + g(2)^2 + \ldots + g(n)^2$
4. **Share of neuron i:** $g(i)^2 \div ( g(1)^2 + g(2)^2 + \ldots + g(n)^2 )$. All the shares add up to 100%
5. **Rank:** sort the neurons by share; the first is "the direction's top neuron"

## WORKED EXAMPLE: QWEN3-8B, LAYER 14, APPLE'S NEURON F7924

| Average on harmful | Average on harmless | Gap | Gap squared | Sum of all 12,288 squared gaps | Share |
|---|---|---|---|---|---|
| −4.41 | −0.16 | −4.24 | 18.0 | about 78 | **18.0 ÷ 78 = 23%** |

> **One neuron out of 12,288 holds 23% of the direction.** If all neurons mattered equally, each would hold 0.008%. No other neuron in the layer holds more; the same is true on Llama (22%) and Qwen3-1.7B (18%).

# 9. The same breakdown in the residual stream

Two readings of the note. A: is any single residual-stream number a big part of the direction? B: how much of the direction does Apple's neuron account for?

## A. THE RESIDUAL DIRECTION OVER ITS OWN NUMBERS

| | Qwen3-8B | Llama-3.1-8B | Qwen3-1.7B |
|---|---|---|---|
| Largest single number's share | 0.8% | 0.5% | 1.0% |
| Top 10 numbers together | 5% | 3% | 7% |
| Spread over about … numbers | 963 of 4,096 | 1,277 of 4,096 | 549 of 2,048 |
| For comparison: largest neuron's share of the MLP direction (slide 16) | 23% | 22% | 18% |

## B. APPLE'S NEURON'S SHARE OF THE RESIDUAL DIRECTION

| | Qwen3-8B | Llama-3.1-8B | Qwen3-1.7B |
|---|---|---|---|
| Layer where Nanda's direction is read | 20 | 12 | 15 |
| Apple's neuron's share of it | 0.5% | 16.9% | 5.3% |
| Its rank among its layer's neurons | 1st of 12,288 | 1st of 14,336 | 1st of 6,144 |
| Largest share of any single neuron | 8.4% (L19:F5693) | 16.9% (Apple's) | 5.3% (Apple's) |

- If A is tiny: no single residual-stream number carries the direction, which would explain why pinning one residual number jailbreaks only 16–32% (slide 16, right)
- If B is small: the residual direction is mostly made by other parts of the model, which would explain why holding the neuron "on" did not undo Nanda's ablation (slide 16, amber row)

**A: yes, no single number matters.** The largest holds 0.5–1.0% of the residual direction, against 18–23% for the top neuron inside the MLP. **B: it depends on the model.** Apple's neuron is 17% of the direction on Llama and 5% on Qwen3-1.7B (the largest of any neuron in both), but 0.5% on Qwen3-8B.

STILL RUNNING LAST TIME: EVERY LANGUAGE

# Pinning the neuron switches refusal off in every language but three

100 HELD-OUT HARMFUL PROMPTS PER LANGUAGE, APPLE'S NEURON AT ITS ENGLISH PIN

| Model | PolyRefuse, 12–13 languages: harmful answers, no pin → pinned | Refused after the pin | MultiJail, 7–8 languages: harmful answers | K.1 criterion, pinned (PolyRefuse / MultiJail) |
|---|---|---|---|---|
| Qwen3-8B | **2–25% → 90–99%** | 0–7% | 2–9% → 66–86% | **95–100% / 79–90% (11 + 8 languages)** |
| Llama-3.1-8B | **4–52% → 87–96%** | 0–6% | 2–29% → 66–78% | *running* |
| Qwen3-1.7B | **14–44% → 83–94%** | 1–14% | 6–20% → 51–76% | *running* |

THE THREE EXCEPTIONS: HARMFUL ANSWERS WITH THE PIN ON

| Language | Qwen3-8B | Llama-3.1-8B | Qwen3-1.7B |
|---|---|---|---|
| **Yoruba** | 1% | 1% | 1% |
| **Swahili** | 3% | 12% | 4% |
| **Javanese** | 36% | 23% | 10% |

In these three the judges also disagree with each other, so the numbers are less certain than elsewhere. Last time this test covered one model and 8 language sets (slide 14).

**One neuron, one pin value, found in English:** it removes refusal in every well-resourced language on all three models. MultiJail's lower harmful rate is its borderline prompts (question 7).

# Inside the MLP, in Arabic, on the two larger models

| | Qwen3-8B, PolyRefuse | Qwen3-8B, MultiJail | Llama, PolyRefuse | Llama, MultiJail |
|---|---|---|---|---|
| Direction's top neuron (its share) | **Apple's (20%)** | Apple's (24%) | **Apple's (10%)** | Apple's (6%) |
| No intervention | 9% | 7% | 13% | 5% |
| Remove the MLP direction | 12% | 9% | 16% | 5% |
| Steer past harmless (α = 4) | **90%** | 69% | **89%** | 21% |
| Pin Apple's neuron | **95%** | 76% | **91%** | 62% |
| Nanda's ablation / with the neuron held on | 85% / 83% | 52% / 49% | 90% / 89% | 59% / 58% |

% = Arabic answers judged harmful (Qwen3Guard). Row 1: the neuron with the biggest share of the Arabic MLP direction. The same table in English is on slide 16.

- Built from Arabic prompts, the direction's biggest neuron is again Apple's neuron
- Removing the direction barely helps on either dataset; pinning the neuron jailbreaks (91–95% on PolyRefuse, 62–76% on MultiJail)
- Holding the neuron on does not undo Nanda's ablation, as in English. Steering is weaker on MultiJail, most of all on Llama (21%)

**All three findings from English hold in Arabic on both larger models.** MultiJail is lower in every row, as everywhere (its borderline prompts, question 7).

# Can the neuron bring refusal back in low-resource languages?

## THE QUESTION

Slide 11: in Yoruba and Swahili the neuron goes quiet and the models refuse less. If we switch the neuron on by hand, does refusal come back?

## HYPOTHESIS

On Llama, yes: its one neuron also creates refusal in English. On the Qwen models probably not, because there the job is shared among several neurons.

## WHAT WE DID

- 100 harmful prompts each in Yoruba, Swahili and Javanese
- Pin the neuron toward its harmful value, at 1×, 2× and 4×
- Count refusals; repeat on 100 harmless prompts to see the cost

## PROMPTS REFUSED: NO PIN → THE TWO STRONGEST PINS

| | Llama: harmful | Llama: harmless | Qwen3-8B: harmful | Qwen3-8B: harmless |
|---|---|---|---|---|
| Yoruba | 11% → 20% / 18% | 5% → 14% / 27% | 62% → 86% / 90% | 34% → 65% / 66% |
| Swahili | **21% → 66% / 73%** | 0% → 18% / 58% | 39% → 89% / 79% | 30% → 76% / 75% |
| Javanese | **61% → 99% / 98%** | **2% → 88% / 100%** | **27% → 96% / 92%** | *pending* |

100 prompts per cell, in that language. Pins: Llama +2 / +4, Qwen3-8B −20 / −40 (toward the neuron's harmful value). Refusal by Qwen3Guard on the original text. Qwen3-1.7B: running.

## SAME IDEA IN ENGLISH, 100 HARMLESS PROMPTS (COMPLETED)

| | Qwen3-8B | Llama-3.1-8B | Qwen3-1.7B |
|---|---|---|---|
| Refused: no pin → strongest pin | 0% → 13% | **0% → 90%** | 0% → 4% |

**The neuron can switch refusal back on (Llama: Swahili and Javanese; Qwen3-8B: Javanese), but it is a blunt switch:** harmless prompts get refused too. Only a mild pin on Llama in Swahili is selective (+1: 43% of harmful, 4% of harmless). In Yoruba, and in Swahili on Qwen3-8B, harmful and harmless rise together.

# Next

## Running, done by tomorrow

- K.1 judge on the per-language attacks, Llama and Qwen3-1.7B
- Rerank in the 9 language sets where Apple's neuron is outside the top 5
- The neuron as one guard across languages, against Llama-Guard and Qwen3Guard
- Arabic residual-stream baseline

## Second halves of the questions

- Q3: overlap for the top 10, 20, 50
- Q5: search and pin on the refused prompts only; a refusal reading that works in Yoruba
- Q6: the layer where the signal first appears, per language
- Q9: the "held on" test at 2× and 4×

## Then

- Job logs in the repository, and a how-to-run guide
- What the neuron responds to (needs a human reader)
- Writing (December), ARR 4 January
- After: images and audio